

# Biswas\_Sayan\_HW2

October 9, 2019

```
[1]: import pandas as pd
import numpy as np
from sklearn import metrics
from sklearn import preprocessing
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
import math
```

## 1 Problem 1

### 1.1 1a

```
[2]: cols = list(pd.read_csv("train.csv", nrows =1))
train = pd.read_csv("train.csv", usecols =[i for i in cols if i not in_
    ↳["id","date","zipcode"]])
train = train.drop(columns=['Unnamed: 0'])

X_train = train.iloc[:,1:]
Y_train = train.iloc[:,0:1]

model_a = LinearRegression().fit(X_train, Y_train)
coeff = pd.DataFrame(model_a.coef_.T,X_train.columns,columns=['Coefficients'])
print("Intercept : ",model_a.intercept_[0])
coeff
```

Intercept : -23088898.60853133

```
[2]:
```

	Coefficients
bedrooms	-14704.280497
bathrooms	25687.783987
sqft_living	83.084210
sqft_lot	0.375930
floors	15555.580988
waterfront	715535.170469
view	63027.898001
condition	18816.402756

grade	79534.602722
sqft_above	42.010495
sqft_basement	41.073715
yr_built	-2400.669330
yr_renovated	43.682942
lat	553505.032276
long	-7424.027121
sqft_living15	68.015792
sqft_lot15	-0.515528

```
[3]: Y_pred = model_a.predict(X_train)
print('Mean Squared Error for train data using model_a:', metrics.
      ↪mean_squared_error(Y_train, Y_pred))
```

Mean Squared Error for train data using model\_a: 31486167775.7949

## 1.2 1b

```
[4]: ss_scaler = preprocessing.StandardScaler()
train_col = train.columns
scaled_data = pd.DataFrame(ss_scaler.fit_transform(train.values), columns =
      ↪train_col)
data = scaled_data.copy()
X_scaled = data.iloc[:,1:]
Y_scaled = data[["price"]]
# model_b = LinearRegression().fit(X_scaled, Y_scaled)

model_b = LinearRegression().fit(X_scaled, Y_scaled)
coeff = pd.DataFrame(model_b.coef_.T, X_scaled.columns, columns=['Coefficients'])
print("Intercept : ", model_b.intercept_[0])
coeff
```

Intercept : 3.086029186623574e-15

```
[4]:
```

	Coefficients
bedrooms	-0.036903
bathrooms	0.054602
sqft_living	0.167243
sqft_lot	0.032070
floors	0.023706
waterfront	0.187856
view	0.142050
condition	0.038207
grade	0.271814
sqft_above	0.142315
sqft_basement	0.079975
yr_built	-0.199350

```

yr_renovated    0.050900
lat             0.230980
long           -0.003051
sqft_living15   0.134321
sqft_lot15     -0.038106

```

```

[5]: Y_pred_scaled = model_b.predict(X_scaled)
      print('Mean Squared Error for training data using model_b:', metrics.
            ↳mean_squared_error(Y_scaled, Y_pred_scaled))

```

Mean Squared Error for training data using model\_b: 0.2734665681293983

### 1.3 1c

```

[6]: cols = list(pd.read_csv("test.csv", nrows =1))
      test = pd.read_csv("test.csv", usecols =[i for i in cols if i not in_
            ↳["id","date","zipcode"]])
      test = test.iloc[:,1:]
      X_test = test.iloc[:,1:]
      Y_test = test.iloc[:,0:1]
      Y_pred_test = model_a.predict(X_test)
      print('Mean Squared Error for test data using model_a:', metrics.
            ↳mean_squared_error(Y_test, Y_pred_test))

```

Mean Squared Error for test data using model\_a: 57628154705.675156

```

[7]: test_col = test.columns
      scaled_data_test = pd.DataFrame(ss_scaler.transform(test.
            ↳values),columns=test_col)
      X_test_scaled = scaled_data_test.iloc[:,1:].values
      Y_test_scaled = scaled_data_test.iloc[:,0:1].values
      Y_pred_test_scaled = model_b.predict(X_test_scaled)
      print('Mean Squared Error for test data using model_b:', metrics.
            ↳mean_squared_error(Y_test_scaled, Y_pred_test_scaled))

```

Mean Squared Error for test data using model\_b: 0.5005173639170148

### 1.4 1d

The coefficients obtained after learning from the unscaled data cannot be compared as they are in different units of measurements respectively. However after feature scaling of the training data, the coefficients can be compared to one another as the features were brought to the same scale before training. The MSE obtained from the unscaled data is a large number and not very intuitive compared to the MSE obtained from the test data to be compared. The features grade, latitude, view, waterfront, sqft\_living, sqft\_living15, sqft\_above and yr\_built mostly contribute to the linear regression model based on the coefficients.

```
[8]: r_squared= metrics.r2_score(Y_test_scaled, Y_pred_test_scaled)
print('coefficient of determination:', r_squared)
```

coefficient of determination: 0.6543560876120955

Based on the R-Squared error, 65.4% of the variation in y can be explained by the dependence on features using the regression model

The model error is 0.5005173639170148.

## 2 Problem 2

### 2.1 2 a,b

```
[9]: def ClosedFormSoln(X, Y):
    if 'x0' not in X.columns:
        x0 = np.ones((X.shape[0], 1), dtype=int)
        X.insert(0, "x0", x0, True)
    X=X.values
    Y=Y.values
    ## calculating the closed form solution to get parameters theta
    X_transpose = np.transpose(X)
    # calculating the dot product
    X_transpose_dotp_x = X_transpose.dot(X)
    # Calculating the inverse
    temp_1 = np.linalg.pinv(X_transpose_dotp_x)

    ##Calculating the second half i.e. (X transpose Y)
    temp_2 = X_transpose.dot(Y)

    ##Calculating theta
    theta = temp_1.dot(temp_2)
    return(theta)

def predict_ClosedFormSoln(theta, X):
    if 'x0' not in X.columns:
        x0 = np.ones((X.shape[0], 1), dtype=int)
        X.insert(0, "x0", x0, True)
    X=X.values
    Y_pred = X.dot(theta)
    return Y_pred
```

```
[10]: #Initializing data
train = scaled_data.copy()
test = scaled_data_test.copy()
# Y values
Y_train = train[["price"]]
Y_test = test[["price"]]
```

```
X_train = train[["sqft_living"]]
X_test = test[["sqft_living"]]
```

```
[11]: # X_test can be a single sample or multiple sample (X_test = 
      ↪test[["sqft_living"]])
X_test = X_test.iloc[0:1,:]
```

```
[12]: theta_single_dim = ClosedFormSoln(X_train,Y_train)
Y_pred = predict_ClosedFormSoln(theta_single_dim,X_test)
print('Response for a new single-dimensional data point: ', Y_pred[0])
```

Response for a new single-dimensional data point: [-0.17565766]

```
[13]: #Initializing data
train = scaled_data.copy()
test = scaled_data_test.copy()
# Y values
Y_train = train[["price"]]
Y_test = test[["price"]]
X_train = train.iloc[:,1:]
X_test = test.iloc[:,1:]
```

```
[14]: # X_test can be a single sample here in the code snippet it takes the first row 
      ↪of the test data
# or multiple rows using (X_test = test.iloc[:,1:])
X_test = X_test.iloc[0:1,:]
```

```
[15]: theta_multi_dim = ClosedFormSoln(X_train,Y_train)
Y_pred = predict_ClosedFormSoln(theta_multi_dim,X_test)
print('Response for a new multi-dimensional data point: ', Y_pred[0])
```

Response for a new multi-dimensional data point: [0.45660638]

## 2.2 2c

Since the closed form gives the global optimum hence the parameters obtained from the closed form solution are same as those obtained from the linear model implemented using the packages. This is the reason the MSE values obtained using the closed form implementation are similar to those obtained using the packages as seen from the output below.

```
[16]: #Initializing data
train = scaled_data.copy()
test = scaled_data_test.copy()
# Y values
Y_train = train[["price"]]
Y_test = test[["price"]]
X_train = train.iloc[:,1:]
```

```
X_test = test.iloc[:,1:]
```

```
[17]: Y_pred_train = predict_ClosedFormSoln(theta_multi_dim,X_train)
Y_pred_test = predict_ClosedFormSoln(theta_multi_dim,X_test)
print('Mean Squared Error for training data using closed form solution:',
      metrics.mean_squared_error(Y_train, Y_pred_train))
print('Mean Squared Error for test data using closed form solution:', metrics.
      mean_squared_error(Y_test, Y_pred_test))
```

Mean Squared Error for training data using closed form solution:

0.2734665681293983

Mean Squared Error for test data using closed form solution: 0.5005173639170145

### 3 Problem 3

#### 3.1 3a

```
[18]: # Compute Cost function is implemented to check if the cost is decreasing with
      iterations
def ComputeCost(X,Y,theta):
    N = Y.shape[0]
    temp = X.dot(theta) - Y
    cost = (1/N)*np.transpose(temp).dot(temp)
    return cost

# Method to compute the theta values using gradient descent
def gradient_descent(X, Y, alpha, num_iters):
    iters = num_iters
    if 'x0' not in X.columns:
        x0 = np.ones((X.shape[0], 1), dtype=int)
        X.insert(0, "x0", x0, True)
    # d is the number of features
    d = X_train.shape[1]
    # Initializing theta with zeros
    theta = np.zeros((d,1))
    X=X.values
    Y=Y.values
    diff_cost = 0
    N = Y.shape[0]
    theta_old = theta
    #while diff_cost > 0.000001 or num_iters > 0
    while num_iters > 0:
        temp = np.transpose(X.dot(theta_old) - Y).dot(X)
        theta_new = theta_old - alpha * (2/N) * np.transpose(temp)
        # checking for convergence
        delta_theta = theta_new - theta_old
        delta = np.sqrt(np.transpose(delta_theta).dot(delta_theta))
```

```

        if delta < 0.00001 and iters != num_iters :
            print("Gradient descent converged at iteration", iters-num_iters)
            break
        num_iters = num_iters - 1
        theta_old = theta_new
        #old_cost = ComputeCost(X,Y,theta_old)
        #new_cost = ComputeCost(X,Y,theta_new)
        #diff_cost = old_cost - new_cost
    return theta_new

# predict the y values using theta and X
def predict_GradientDescent(theta, X):
    if 'x0' not in X.columns:
        x0 = np.ones((X.shape[0], 1), dtype=int)
        X.insert(0, "x0", x0, True)
    X = X.values
    Y_pred = X.dot(theta)
    return Y_pred

```

### 3.2 3b

```

[19]: train = scaled_data.copy()
      test = scaled_data_test.copy()
      # Y values
      Y_train = train[["price"]]
      Y_test = test[["price"]]
      X_train = train.iloc[:,1:]
      X_test = test.iloc[:,1:]

[20]: alpha = [0.01, 0.03, 0.1]
      num_iters = [10, 50, 100]

      for i in alpha:
          for j in num_iters :
              theta = gradient_descent(X_train, Y_train, i, j)
              print('theta when alpha = ' + str(i) + ' and num_iters = ' + str(j) + ':
→\n', theta)
              print("\n")
              Y_pred_train = predict_GradientDescent(theta, X_train)
              Y_pred_test = predict_GradientDescent(theta, X_test)
              print('MSE for training data using Gradient descent when alpha = ' +
→str(i) +
                  ' and num_iters = ' + str(j) + ':', metrics.
→mean_squared_error(Y_train, Y_pred_train))
              print('MSE for testing data using Gradient descent when alpha = ' +
→str(i) +

```

```

        ' and num_iters = '+ str(j) + ':', metrics.
    →mean_squared_error(Y_test, Y_pred_test))
    print("\n")

```

theta when alpha = 0.01 and num\_iters = 10:

```

[[ 8.24229573e-17]
 [ 3.51546174e-02]
 [ 5.97764565e-02]
 [ 9.46499119e-02]
 [ 1.58577098e-02]
 [ 2.80724205e-02]
 [ 5.14216548e-02]
 [ 6.74806965e-02]
 [ 1.46048174e-02]
 [ 8.81429767e-02]
 [ 7.58100530e-02]
 [ 5.35522376e-02]
 [-1.19508461e-02]
 [ 2.46116923e-02]
 [ 6.15128542e-02]
 [-4.24986932e-03]
 [ 8.66260205e-02]
 [ 1.73904498e-02]]

```

MSE for training data using Gradient descent when alpha = 0.01 and num\_iters = 10: 0.4769826562643797

MSE for testing data using Gradient descent when alpha = 0.01 and num\_iters = 10: 0.7899805878552555

theta when alpha = 0.01 and num\_iters = 50:

```

[[ 6.05232531e-16]
 [ 1.76980518e-02]
 [ 7.02037279e-02]
 [ 1.61104777e-01]
 [ 1.07199113e-02]
 [ 3.27773917e-02]
 [ 1.35936441e-01]
 [ 1.46319888e-01]
 [ 4.29116775e-02]
 [ 1.66822006e-01]
 [ 1.25277756e-01]
 [ 9.77393699e-02]
 [-9.43446724e-02]
 [ 6.27487154e-02]

```



```
[ 1.77221329e-01]
[-3.88336062e-02]
[ 1.51099990e-01]
[ 7.90555822e-03]]
```

```
MSE for training data using Gradient descent when alpha = 0.01 and num_iters =
50: 0.2935883463907345
MSE for testing data using Gradient descent when alpha = 0.01 and num_iters =
50: 0.5343327522759017
```

```
theta when alpha = 0.01 and num_iters = 100:
```

```
[[ 1.23490662e-15]
[-1.07772411e-02]
[ 5.67070436e-02]
[ 1.67851095e-01]
[ 7.93841518e-03]
[ 3.07915991e-02]
[ 1.65358934e-01]
[ 1.54682331e-01]
[ 4.72356963e-02]
[ 1.95563835e-01]
[ 1.37140433e-01]
[ 9.02383856e-02]
[-1.35234073e-01]
[ 6.71310263e-02]
[ 2.21279091e-01]
[-3.74686822e-02]
[ 1.61148092e-01]
[-3.15262960e-03]]
```

```
MSE for training data using Gradient descent when alpha = 0.01 and num_iters =
100: 0.27842422313332665
MSE for testing data using Gradient descent when alpha = 0.01 and num_iters =
100: 0.5124065597546549
```

```
theta when alpha = 0.03 and num_iters = 10:
```

```
[[ 3.59918761e-16]
[ 3.44783187e-02]
[ 7.77279035e-02]
[ 1.51030380e-01]
[ 1.49096249e-02]
[ 3.45824114e-02]
[ 1.09451543e-01]
[ 1.28344010e-01]
```

```
[ 3.53335842e-02]
[ 1.47384928e-01]
[ 1.16499005e-01]
[ 9.32828025e-02]
[-6.15572702e-02]
[ 5.28254858e-02]
[ 1.38219844e-01]
[-2.88618577e-02]
[ 1.39434856e-01]
[ 1.51205380e-02]]
```

MSE for training data using Gradient descent when alpha = 0.03 and num\_iters = 10: 0.32025061817365497  
MSE for testing data using Gradient descent when alpha = 0.03 and num\_iters = 10: 0.5706630290510314

theta when alpha = 0.03 and num\_iters = 50:

```
[[ 1.81494153e-15]
[-2.42677316e-02]
[ 5.09297569e-02]
[ 1.69308284e-01]
[ 8.93562133e-03]
[ 2.85160307e-02]
[ 1.75952382e-01]
[ 1.50939455e-01]
[ 4.68466787e-02]
[ 2.13358145e-01]
[ 1.42588380e-01]
[ 8.35618907e-02]
[-1.54222591e-01]
[ 6.42322034e-02]
[ 2.34214776e-01]
[-2.88323918e-02]
[ 1.62387196e-01]
[-1.01317941e-02]]
```

MSE for training data using Gradient descent when alpha = 0.03 and num\_iters = 50: 0.27560896108344324  
MSE for testing data using Gradient descent when alpha = 0.03 and num\_iters = 50: 0.5074745831005474

theta when alpha = 0.03 and num\_iters = 100:

```
[[ 2.66659361e-15]
[-3.47525840e-02]
```

```
[ 4.93272743e-02]
[ 1.69715918e-01]
[ 1.56037395e-02]
[ 2.41618518e-02]
[ 1.84600950e-01]
[ 1.43066441e-01]
[ 4.33028607e-02]
[ 2.41936788e-01]
[ 1.45193915e-01]
[ 7.97991240e-02]
[-1.79016692e-01]
[ 5.72099614e-02]
[ 2.35964039e-01]
[-1.42836817e-02]
[ 1.54438690e-01]
[-2.21623928e-02]]
```

MSE for training data using Gradient descent when  $\alpha = 0.03$  and  $\text{num\_iters} = 100$ : 0.27395671526270826

MSE for testing data using Gradient descent when  $\alpha = 0.03$  and  $\text{num\_iters} = 100$ : 0.5030789048604056

theta when  $\alpha = 0.1$  and  $\text{num\_iters} = 10$ :

```
[[ 1.24420474e-15]
[-1.36910568e-02]
[ 5.52088472e-02]
[ 1.68353541e-01]
[ 7.45623778e-03]
[ 3.10991810e-02]
[ 1.68132153e-01]
[ 1.55868304e-01]
[ 4.76371951e-02]
[ 1.97351136e-01]
[ 1.38347012e-01]
[ 8.91136198e-02]
[-1.38420055e-01]
[ 6.78915154e-02]
[ 2.26265130e-01]
[-3.78793254e-02]
[ 1.62485721e-01]
[-3.72317254e-03]]
```

MSE for training data using Gradient descent when  $\alpha = 0.1$  and  $\text{num\_iters} = 10$ : 0.2777784544544066

MSE for testing data using Gradient descent when  $\alpha = 0.1$  and  $\text{num\_iters} = 10$ :

0.5113613440792656

theta when alpha = 0.1 and num\_iters = 50:

```
[[ 2.94542168e-15]
 [-3.67123703e-02]
 [ 5.19311794e-02]
 [ 1.68537088e-01]
 [ 2.31442397e-02]
 [ 2.35943958e-02]
 [ 1.86936278e-01]
 [ 1.41325823e-01]
 [ 4.01410188e-02]
 [ 2.59030114e-01]
 [ 1.43764344e-01]
 [ 7.99826561e-02]
 [-1.91538284e-01]
 [ 5.33598604e-02]
 [ 2.32971158e-01]
 [-7.57832894e-03]
 [ 1.44321559e-01]
 [-2.98970970e-02]]
```

MSE for training data using Gradient descent when alpha = 0.1 and num\_iters = 50: 0.27356291763134366

MSE for testing data using Gradient descent when alpha = 0.1 and num\_iters = 50: 0.5013370393181915

theta when alpha = 0.1 and num\_iters = 100:

```
[[ 3.07953663e-15]
 [-3.68947451e-02]
 [ 5.42975971e-02]
 [ 1.67304747e-01]
 [ 3.02686805e-02]
 [ 2.37564001e-02]
 [ 1.87723494e-01]
 [ 1.41848046e-01]
 [ 3.83994393e-02]
 [ 2.70127986e-01]
 [ 1.42361073e-01]
 [ 8.00147289e-02]
 [-1.98462669e-01]
 [ 5.11892697e-02]
 [ 2.31198797e-01]
 [-3.65404153e-03]
 [ 1.35919784e-01]]
```

```
[-3.63918536e-02]]
```

```
MSE for training data using Gradient descent when alpha = 0.1 and num_iters =  
100: 0.27346886118648905  
MSE for testing data using Gradient descent when alpha = 0.1 and num_iters =  
100: 0.5006111869631489
```

### 3.3 3c

The MSE value for the training and the testing data using Gradient Descent for Alpha = 0.1 and iterations = 100 with theta being initialized with all zeroes; is almost similar to those obtained with the package. The more the number of iterations, the closer the MSE is to the one obtained using the package. We can observe that even with smaller learning rate, the algorithm converges if provided with sufficient number of iterations. The MSE value is small for a bigger learning rate compared to the small learning rates for the same number of iterations. The objective decreases for the selected learning rates with each iterations however if the learning rate is selected as 0.2 the cost increases with each iteration and hence cannot be selected as an alpha value for gradient descent. Using the above implementation, the algorithm is converging with the number of iterations reaching the maximum number of iterations.

### 3.4 3d

```
[21]: # Compute Cost function is implemented to check if the cost is decreasing with
      ↪ iterations
def ComputeCost(X,Y,theta):
    N = Y.shape[0]
    temp = X.dot(theta) - Y
    cost = (1/N)*np.transpose(temp).dot(temp)
    return cost

# Method to compute the theta values using gradient descent
def gradient_descent_line_search(X, Y, alpha_max, step, e, T):
    if 'x0' not in X.columns:
        x0 = np.ones((X.shape[0], 1), dtype=int)
        X.insert(0, "x0", x0, True)
    # d is the number of features
    d = X_train.shape[1]
    # Initializing theta with zeros
    theta = np.zeros((d,1))

    X=X.values
    Y=Y.values
    N = Y.shape[0]
    delta = 1
    alpha = alpha_max
    backtrack = T
    while(delta > 0.0001):
        T = backtrack
        theta_try = theta
        while(T>0):
            temp = np.transpose(X.dot(theta) - Y).dot(X)
            theta_try = theta - alpha * (2/N) * np.transpose(temp)
            if abs((ComputeCost(X,Y, theta) - ComputeCost(X,Y, theta_try))) > e:
                theta = theta_try
                #break
            else:
                alpha = step * alpha
                T = T-1
            delta_theta = theta_try - theta
            delta = np.sqrt(np.transpose(delta_theta).dot(delta_theta))
        return theta

# predict the y values using theta and X
def predict_GradientDescent(theta, X):
    if 'x0' not in X.columns:
        x0 = np.ones((X.shape[0], 1), dtype=int)
```

```

        X.insert(0, "x0", x0, True)
    X = X.values
    Y_pred = X.dot(theta)
    return Y_pred

```

```

[22]: train = scaled_data.copy()
test = scaled_data_test.copy()
# Y values
Y_train = train[["price"]]
Y_test = test[["price"]]
X_train = train.iloc[:,1:]
X_test = test.iloc[:,1:]

```

```

[23]: alpha = [0.01, 0.03, 0.1]
backtrack = [10, 50, 100]

for i in alpha:
    for j in backtrack :
        theta = gradient_descent_line_search(X_train, Y_train, i, 0.5, 0.0001, j)
        #print('theta when alpha = ' + str(i) + ' and num_iters = ' + str(j) + ':
→\n', theta)
        #print("\n")
        Y_pred_train = predict_GradientDescent(theta, X_train)
        Y_pred_test = predict_GradientDescent(theta, X_test)
        print('MSE for training data using Gradient descent when alpha = ' +
→str(i) +
        ' and backtrack = ' + str(j) + ':', metrics.
→mean_squared_error(Y_train, Y_pred_train))
        print('MSE for testing data using Gradient descent when alpha = ' +
→str(i) +
        ' and backtrack = ' + str(j) + ':', metrics.
→mean_squared_error(Y_test, Y_pred_test))
        print("\n")

```

```

MSE for training data using Gradient descent when alpha = 0.01 and backtrack =
10: 0.4769826562643797
MSE for testing data using Gradient descent when alpha = 0.01 and backtrack =
10: 0.7899805878552555

```

```

MSE for training data using Gradient descent when alpha = 0.01 and backtrack =
50: 0.2935883463907345
MSE for testing data using Gradient descent when alpha = 0.01 and backtrack =
50: 0.5343327522759017

```

```

MSE for training data using Gradient descent when alpha = 0.01 and backtrack =

```

```

100: 0.27842422313332665
MSE for testing data using Gradient descent when alpha = 0.01 and backtrack =
100: 0.5124065597546549

MSE for training data using Gradient descent when alpha = 0.03 and backtrack =
10: 0.32025061817365497
MSE for testing data using Gradient descent when alpha = 0.03 and backtrack =
10: 0.5706630290510314

MSE for training data using Gradient descent when alpha = 0.03 and backtrack =
50: 0.27588674681517505
MSE for testing data using Gradient descent when alpha = 0.03 and backtrack =
50: 0.5080309707498859

MSE for training data using Gradient descent when alpha = 0.03 and backtrack =
100: 0.27588674681517505
MSE for testing data using Gradient descent when alpha = 0.03 and backtrack =
100: 0.5080309707498859

MSE for training data using Gradient descent when alpha = 0.1 and backtrack =
10: 0.2777784544544066
MSE for testing data using Gradient descent when alpha = 0.1 and backtrack = 10:
0.5113613440792656

MSE for training data using Gradient descent when alpha = 0.1 and backtrack =
50: 0.2745099563397425
MSE for testing data using Gradient descent when alpha = 0.1 and backtrack = 50:
0.5048703004874806

MSE for training data using Gradient descent when alpha = 0.1 and backtrack =
100: 0.2745099563397425
MSE for testing data using Gradient descent when alpha = 0.1 and backtrack =
100: 0.5048703004874806

```

The values of MSE in the training and testing data using gradient descent using line search is almost similar to those obtained using Gradient descent.

## 4 Problem 4



## Problem 4

$$a) J(\theta) = \frac{1}{2} \sum_{i=1}^N (h_{\theta}(x_i) - y_i)^2 + \frac{1}{2} \lambda \sum_{j=1}^d \theta_j^2$$

$J(\theta)$  can also be written as, the below equation is vector notation.

$$J(\theta) = \frac{1}{2} \|X\theta - y\|^2 + \frac{\lambda}{2} \|\theta\|^2$$

To find,  
 $\min_{\theta} J(\theta),$

$$\begin{aligned} \frac{\partial J(\theta)}{\partial \theta} &= \frac{1}{2} \times 2 (X\theta - y)^T X + \frac{\lambda}{2} \times 2 \theta^T \\ &= (X\theta - y)^T X + \lambda \theta^T \end{aligned}$$

Equating  $\frac{\partial J(\theta)}{\partial \theta} = 0$  to find the minimum,

$$\Rightarrow (X\theta - y)^T X + \lambda \theta^T = 0$$

Transposing on both the sides, we get,

$$\Rightarrow ((X\theta - y)^T X + \lambda \theta^T)^T = 0^T$$

$$\Rightarrow ((X\theta - y)^T X)^T + (\lambda \theta^T)^T = 0$$

$$\Rightarrow X^T (X\theta - y) + \lambda \theta = 0$$

$$\Rightarrow X^T X \theta - X^T y + \lambda \theta = 0$$

$$\Rightarrow X^T X \theta + \lambda \theta - X^T y = 0$$

$$\Rightarrow (X^T X + \lambda I) \theta - X^T y = 0$$

$$[\because \lambda I = \lambda]$$

$$\Rightarrow \theta = (X^T X + \lambda I)^{-1} X^T y$$

## 4.1 4b

```
[24]: def ClosedForm_RidgeRegression(X, Y, lambda_r):  
    if 'x0' not in X.columns:  
        x0 = np.ones((X.shape[0], 1), dtype=int)  
        X.insert(0, "x0", x0, True)  
    X=X.values  
    Y=Y.values  
    ## calculating the closed form solution to get parameters theta  
    X_transpose = np.transpose(X)  
    # calculating the dot product  
    X_transpose_dotp_x = X_transpose.dot(X)  
  
    # Matrix for Regularization term  
    L = np.identity(X.shape[1])  
    #No regularization/penalizing on the theta0  
    L[0][0]=0  
    # Calculating the inverse  
    temp_1 = np.linalg.pinv(X_transpose_dotp_x + (lambda_r * L))  
  
    ##Calculating the second half i.e. (X transpose Y)  
    temp_2 = X_transpose.dot(Y)  
  
    ##Calculating theta  
    theta = temp_1.dot(temp_2)  
    return(theta)  
  
def predict(theta, X):  
    if 'x0' not in X.columns:  
        x0 = np.ones((X.shape[0], 1), dtype=int)  
        X.insert(0, "x0", x0, True)  
    Y_pred = X.dot(theta)  
    return Y_pred
```

```
[25]: #Initializing data  
train = scaled_data.copy()  
test = scaled_data_test.copy()  
# Y values  
Y_train = train[["price"]]  
Y_test = test[["price"]]  
X_train = train.iloc[:,1:]  
X_test = test.iloc[:,1:]
```

```
[26]: #lambda_list = [0,1,2,5,10,15,20,30,50,75,100,500,1000,2000,3000,5000]  
#lambda_list = [0,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1]  
lambda_list = []  
MSE_train = []
```

```

#for i in np.arange(1,100,1):
for i in range(1,101,1):
    #for i in lambda_list:
        lambda_list.append(i)
        theta = ClosedFormRidgeRegression(X_train,Y_train,i)
        #print(theta)
        Y_pred = predict(theta,X_train)
        print('Mean Squared Error for training data using closed form solution with_λ
→lambda= ' + str(i) + ': ',
              metrics.mean_squared_error(Y_train, Y_pred))
        #print("\n")
        MSE_train.append(metrics.mean_squared_error(Y_train, Y_pred))

min_lambda = lambda_list[MSE_train.index(min(MSE_train))]
print("\n")
print("Minimum MSE value for the training data = ", min(MSE_train), "obtained_λ
→for lambda =", min_lambda)
plt.plot(lambda_list,MSE_train)
plt.title("Lambda v/s MSE on the training data ")
plt.xlabel("Lambda")
plt.ylabel("MSE")
plt.show()

```

```

Mean Squared Error for training data using closed form solution with lambda= 1:
0.2734668788542625
Mean Squared Error for training data using closed form solution with lambda= 2:
0.27346780421401556
Mean Squared Error for training data using closed form solution with lambda= 3:
0.2734693341886705
Mean Squared Error for training data using closed form solution with lambda= 4:
0.2734714590206986
Mean Squared Error for training data using closed form solution with lambda= 5:
0.2734741692062795
Mean Squared Error for training data using closed form solution with lambda= 6:
0.27347745548689273
Mean Squared Error for training data using closed form solution with lambda= 7:
0.2734813088412359
Mean Squared Error for training data using closed form solution with lambda= 8:
0.2734857204774564
Mean Squared Error for training data using closed form solution with lambda= 9:
0.2734906818256808
Mean Squared Error for training data using closed form solution with lambda= 10:
0.2734961845308304
Mean Squared Error for training data using closed form solution with lambda= 11:
0.2735022204457091
Mean Squared Error for training data using closed form solution with lambda= 12:
0.2735087816243536

```

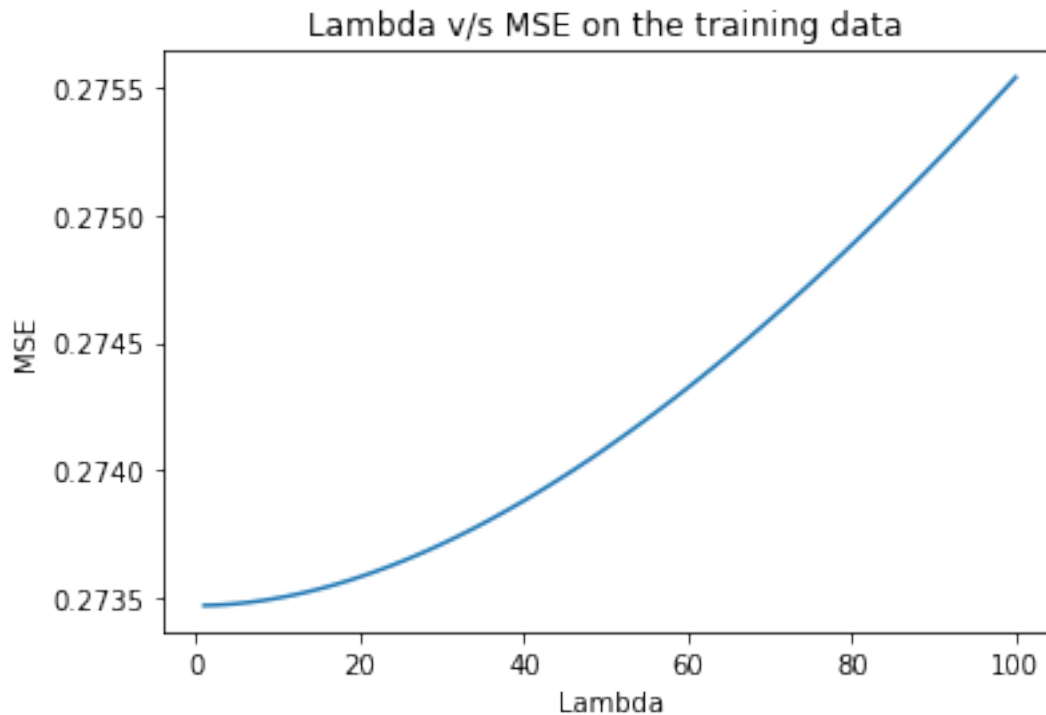
Mean Squared Error for training data using closed form solution with  $\lambda = 13$ :  
0.27351586031563274  
Mean Squared Error for training data using closed form solution with  $\lambda = 14$ :  
0.2735234489570864  
Mean Squared Error for training data using closed form solution with  $\lambda = 15$ :  
0.27353154016899406  
Mean Squared Error for training data using closed form solution with  $\lambda = 16$ :  
0.27354012674866146  
Mean Squared Error for training data using closed form solution with  $\lambda = 17$ :  
0.27354920166491864  
Mean Squared Error for training data using closed form solution with  $\lambda = 18$ :  
0.2735587580528187  
Mean Squared Error for training data using closed form solution with  $\lambda = 19$ :  
0.27356878920852923  
Mean Squared Error for training data using closed form solution with  $\lambda = 20$ :  
0.27357928858440955  
Mean Squared Error for training data using closed form solution with  $\lambda = 21$ :  
0.27359024978426405  
Mean Squared Error for training data using closed form solution with  $\lambda = 22$ :  
0.2736016665587666  
Mean Squared Error for training data using closed form solution with  $\lambda = 23$ :  
0.27361353280104744  
Mean Squared Error for training data using closed form solution with  $\lambda = 24$ :  
0.2736258425424361  
Mean Squared Error for training data using closed form solution with  $\lambda = 25$ :  
0.27363858994835594  
Mean Squared Error for training data using closed form solution with  $\lambda = 26$ :  
0.2736517693143609  
Mean Squared Error for training data using closed form solution with  $\lambda = 27$ :  
0.2736653750623118  
Mean Squared Error for training data using closed form solution with  $\lambda = 28$ :  
0.27367940173668504  
Mean Squared Error for training data using closed form solution with  $\lambda = 29$ :  
0.2736938440010089  
Mean Squared Error for training data using closed form solution with  $\lambda = 30$ :  
0.27370869663442243  
Mean Squared Error for training data using closed form solution with  $\lambda = 31$ :  
0.2737239545283518  
Mean Squared Error for training data using closed form solution with  $\lambda = 32$ :  
0.27373961268330016  
Mean Squared Error for training data using closed form solution with  $\lambda = 33$ :  
0.2737556662057459  
Mean Squared Error for training data using closed form solution with  $\lambda = 34$ :  
0.2737721103051456  
Mean Squared Error for training data using closed form solution with  $\lambda = 35$ :  
0.2737889402910377  
Mean Squared Error for training data using closed form solution with  $\lambda = 36$ :  
0.2738061515702425

Mean Squared Error for training data using closed form solution with  $\lambda = 37$ :  
0.2738237396441549  
Mean Squared Error for training data using closed form solution with  $\lambda = 38$ :  
0.27384170010612796  
Mean Squared Error for training data using closed form solution with  $\lambda = 39$ :  
0.2738600286389406  
Mean Squared Error for training data using closed form solution with  $\lambda = 40$ :  
0.2738787210123495  
Mean Squared Error for training data using closed form solution with  $\lambda = 41$ :  
0.273897773080721  
Mean Squared Error for training data using closed form solution with  $\lambda = 42$ :  
0.2739171807807383  
Mean Squared Error for training data using closed form solution with  $\lambda = 43$ :  
0.273936940129184  
Mean Squared Error for training data using closed form solution with  $\lambda = 44$ :  
0.27395704722079367  
Mean Squared Error for training data using closed form solution with  $\lambda = 45$ :  
0.27397749822617773  
Mean Squared Error for training data using closed form solution with  $\lambda = 46$ :  
0.2739982893898091  
Mean Squared Error for training data using closed form solution with  $\lambda = 47$ :  
0.2740194170280754  
Mean Squared Error for training data using closed form solution with  $\lambda = 48$ :  
0.27404087752739237  
Mean Squared Error for training data using closed form solution with  $\lambda = 49$ :  
0.27406266734237567  
Mean Squared Error for training data using closed form solution with  $\lambda = 50$ :  
0.27408478299407113  
Mean Squared Error for training data using closed form solution with  $\lambda = 51$ :  
0.2741072210682389  
Mean Squared Error for training data using closed form solution with  $\lambda = 52$ :  
0.27412997821369106  
Mean Squared Error for training data using closed form solution with  $\lambda = 53$ :  
0.27415305114068067  
Mean Squared Error for training data using closed form solution with  $\lambda = 54$ :  
0.27417643661934016  
Mean Squared Error for training data using closed form solution with  $\lambda = 55$ :  
0.2742001314781671  
Mean Squared Error for training data using closed form solution with  $\lambda = 56$ :  
0.2742241326025562  
Mean Squared Error for training data using closed form solution with  $\lambda = 57$ :  
0.27424843693337575  
Mean Squared Error for training data using closed form solution with  $\lambda = 58$ :  
0.27427304146558695  
Mean Squared Error for training data using closed form solution with  $\lambda = 59$ :  
0.2742979432469046  
Mean Squared Error for training data using closed form solution with  $\lambda = 60$ :  
0.2743231393764978

Mean Squared Error for training data using closed form solution with  $\lambda = 61$ :  
 0.2743486270037289  
 Mean Squared Error for training data using closed form solution with  $\lambda = 62$ :  
 0.2743744033269306  
 Mean Squared Error for training data using closed form solution with  $\lambda = 63$ :  
 0.2744004655922183  
 Mean Squared Error for training data using closed form solution with  $\lambda = 64$ :  
 0.2744268110923376  
 Mean Squared Error for training data using closed form solution with  $\lambda = 65$ :  
 0.2744534371655453  
 Mean Squared Error for training data using closed form solution with  $\lambda = 66$ :  
 0.27448034119452336  
 Mean Squared Error for training data using closed form solution with  $\lambda = 67$ :  
 0.27450752060532385  
 Mean Squared Error for training data using closed form solution with  $\lambda = 68$ :  
 0.274534972866344  
 Mean Squared Error for training data using closed form solution with  $\lambda = 69$ :  
 0.27456269548733175  
 Mean Squared Error for training data using closed form solution with  $\lambda = 70$ :  
 0.27459068601841896  
 Mean Squared Error for training data using closed form solution with  $\lambda = 71$ :  
 0.2746189420491818  
 Mean Squared Error for training data using closed form solution with  $\lambda = 72$ :  
 0.2746474612077288  
 Mean Squared Error for training data using closed form solution with  $\lambda = 73$ :  
 0.2746762411598135  
 Mean Squared Error for training data using closed form solution with  $\lambda = 74$ :  
 0.2747052796079727  
 Mean Squared Error for training data using closed form solution with  $\lambda = 75$ :  
 0.27473457429068865  
 Mean Squared Error for training data using closed form solution with  $\lambda = 76$ :  
 0.27476412298157504  
 Mean Squared Error for training data using closed form solution with  $\lambda = 77$ :  
 0.27479392348858406  
 Mean Squared Error for training data using closed form solution with  $\lambda = 78$ :  
 0.27482397365323746  
 Mean Squared Error for training data using closed form solution with  $\lambda = 79$ :  
 0.274854271349877  
 Mean Squared Error for training data using closed form solution with  $\lambda = 80$ :  
 0.2748848144849361  
 Mean Squared Error for training data using closed form solution with  $\lambda = 81$ :  
 0.2749156009962318  
 Mean Squared Error for training data using closed form solution with  $\lambda = 82$ :  
 0.274946628852275  
 Mean Squared Error for training data using closed form solution with  $\lambda = 83$ :  
 0.27497789605160017  
 Mean Squared Error for training data using closed form solution with  $\lambda = 84$ :  
 0.2750094006221126

Mean Squared Error for training data using closed form solution with  $\lambda = 85$ :  
0.27504114062045326  
Mean Squared Error for training data using closed form solution with  $\lambda = 86$ :  
0.27507311413138086  
Mean Squared Error for training data using closed form solution with  $\lambda = 87$ :  
0.27510531926716997  
Mean Squared Error for training data using closed form solution with  $\lambda = 88$ :  
0.27513775416702535  
Mean Squared Error for training data using closed form solution with  $\lambda = 89$ :  
0.27517041699651146  
Mean Squared Error for training data using closed form solution with  $\lambda = 90$ :  
0.27520330594699766  
Mean Squared Error for training data using closed form solution with  $\lambda = 91$ :  
0.27523641923511644  
Mean Squared Error for training data using closed form solution with  $\lambda = 92$ :  
0.27526975510223745  
Mean Squared Error for training data using closed form solution with  $\lambda = 93$ :  
0.2753033118139544  
Mean Squared Error for training data using closed form solution with  $\lambda = 94$ :  
0.2753370876595843  
Mean Squared Error for training data using closed form solution with  $\lambda = 95$ :  
0.27537108095168217  
Mean Squared Error for training data using closed form solution with  $\lambda = 96$ :  
0.27540529002556463  
Mean Squared Error for training data using closed form solution with  $\lambda = 97$ :  
0.27543971323884886  
Mean Squared Error for training data using closed form solution with  $\lambda = 98$ :  
0.27547434897100154  
Mean Squared Error for training data using closed form solution with  $\lambda = 99$ :  
0.27550919562289927  
Mean Squared Error for training data using closed form solution with  $\lambda = 100$ : 0.2755442516164008

Minimum MSE value for the training data = 0.2734668788542625 obtained for  $\lambda = 1$



```
[27]: MSE_test = []
lambda_list = []
#for i in np.arange(1e-5,10.0,1e-2):
for i in range(1,101,1):
    lambda_list.append(i)
    theta = ClosedForm_RidgeRegression(X_train,Y_train,i)
    Y_pred = predict(theta,X_test)
    print('Mean Squared Error for testing data using closed form solution with_
    →lambda= ' + str(i) + ': ',
          metrics.mean_squared_error(Y_test, Y_pred))
    MSE_test.append(metrics.mean_squared_error(Y_test, Y_pred))
print("\n")
min_lambda = lambda_list[MSE_test.index(min(MSE_test))]
print("Minimum MSE value for the testing data = ", min(MSE_test), "obtained for_
    →lambda =", min_lambda)
```

Mean Squared Error for testing data using closed form solution with lambda= 1:  
0.5005925887198549

Mean Squared Error for testing data using closed form solution with lambda= 2:  
0.5006683305125841

Mean Squared Error for testing data using closed form solution with lambda= 3:



0.5007445752685044  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 4$ :  
 0.5008213094525349  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 5$ :  
 0.500898520001577  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 6$ :  
 0.5009761943057558  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 7$ :  
 0.5010543201904955  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 8$ :  
 0.501132885899432  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 9$ :  
 0.50121188007802  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 10$ :  
 0.5012912917578864  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 11$ :  
 0.5013711103418312  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 12$ :  
 0.5014513255895406  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 13$ :  
 0.5015319276038572  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 14$ :  
 0.5016129068176468  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 15$ :  
 0.5016942539812503  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 16$ :  
 0.5017759601504561  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 17$ :  
 0.5018580166749654  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 18$ :  
 0.5019404151873621  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 19$ :  
 0.5020231475925323  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 20$ :  
 0.502106206057529  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 21$ :  
 0.5021895830018502  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 22$ :  
 0.5022732710881335  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 23$ :  
 0.5023572632132074  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 24$ :  
 0.5024415524995294  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 25$ :  
 0.5025261322869632  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 26$ :  
 0.5026109961248924  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 27$ :

0.5026961377646414  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 28$ :  
 0.5027815511522193  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 29$ :  
 0.5028672304213353  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 30$ :  
 0.5029531698866975  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 31$ :  
 0.5030393640375841  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 32$ :  
 0.5031258075316567  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 33$ :  
 0.5032124951890279  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 34$ :  
 0.5032994219865451  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 35$ :  
 0.5033865830523175  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 36$ :  
 0.5034739736604354  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 37$ :  
 0.5035615892259011  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 38$ :  
 0.5036494252997626  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 39$ :  
 0.5037374775644202  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 40$ :  
 0.5038257418291154  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 41$ :  
 0.503914214025602  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 42$ :  
 0.5040028902039613  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 43$ :  
 0.5040917665285937  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 44$ :  
 0.5041808392743496  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 45$ :  
 0.5042701048228084  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 46$ :  
 0.5043595596586953  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 47$ :  
 0.5044492003664289  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 48$ :  
 0.5045390236268023  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 49$ :  
 0.5046290262137791  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 50$ :  
 0.5047192049914125  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 51$ :

0.5048095569108725  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 52$ :  
 0.5049000790075835  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 53$ :  
 0.5049907683984668  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 54$ :  
 0.5050816222792798  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 55$ :  
 0.5051726379220536  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 56$ :  
 0.5052638126726183  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 57$ :  
 0.5053551439482235  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 58$ :  
 0.5054466292352361  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 59$ :  
 0.5055382660869254  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 60$ :  
 0.5056300521213241  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 61$ :  
 0.5057219850191623  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 62$ :  
 0.5058140625218792  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 63$ :  
 0.5059062824297016  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 64$ :  
 0.5059986425997876  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 65$ :  
 0.5060911409444377  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 66$ :  
 0.5061837754293654  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 67$ :  
 0.5062765440720337  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 68$ :  
 0.506369444940034  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 69$ :  
 0.5064624761495409  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 70$ :  
 0.5065556358638036  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 71$ :  
 0.5066489222916962  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 72$ :  
 0.5067423336863143  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 73$ :  
 0.5068358683436214  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 74$ :  
 0.5069295246011394  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 75$ :

0.5070233008366813  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 76$ :  
 0.5071171954671343  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 77$ :  
 0.5072112069472702  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 78$ :  
 0.5073053337686099  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 79$ :  
 0.5073995744583145  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 80$ :  
 0.5074939275781184  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 81$ :  
 0.5075883917232973  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 82$ :  
 0.5076829655216698  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 83$ :  
 0.5077776476326278  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 84$ :  
 0.5078724367462084  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 85$ :  
 0.5079673315821837  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 86$ :  
 0.5080623308891917  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 87$ :  
 0.5081574334438861  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 88$ :  
 0.5082526380501177  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 89$ :  
 0.5083479435381452  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 90$ :  
 0.5084433487638663  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 91$ :  
 0.508538852608074  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 92$ :  
 0.5086344539757397  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 93$ :  
 0.5087301517953179  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 94$ :  
 0.5088259450180703  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 95$ :  
 0.5089218326174174  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 96$ :  
 0.5090178135883017  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 97$ :  
 0.5091138869465809  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 98$ :  
 0.5092100517284336  
 Mean Squared Error for testing data using closed form solution with  $\lambda = 99$ :

0.5093063069897821


Mean Squared Error for testing data using closed form solution with  $\lambda = 100$ :

0.5094026518057427

Minimum MSE value for the testing data = 0.5005925887198549 obtained for  $\lambda = 1$

```
[28]: def annot_min(x,y, ax=None):
    xmin = x[np.argmin(y)]
    ymin = min(y)
    text= "x={:.3f}, y={:.3f}".format(xmin, ymin)
    if not ax:
        ax=plt.gca()
    bbox_props = dict(boxstyle="square,pad=0.3", fc="w", ec="k", lw=0.72)
    arrowprops=dict(arrowstyle="->",connectionstyle="angle,angleA=0,angleB=60")
    kw = dict(xycoords='data',textcoords="axes fraction",
              arrowprops=arrowprops, bbox=bbox_props, ha="right", va="top")
    ax.annotate(text, xy=(xmin, ymin), xytext=(0.50,0.50), **kw)

plt.plot(lambda_list, MSE_test)
annot_min(lambda_list, MSE_test)
plt.title("Lambda v/s MSE on the testing data")
plt.xlabel("Lambda")
plt.ylabel("MSE")
plt.show()
```



The MSE value for the training data set obtained in the Ridge regression is almost similar to the value obtained in the Linear Regression. The MSE value for the testing data set obtained in the Ridge regression is slightly less compared to those obtained in the Linear Regression.

# Problem 5

$$X = \begin{bmatrix} 1 & x_{11} & x_{12} \\ 1 & x_{21} & x_{22} \\ \vdots & \vdots & \vdots \\ 1 & x_{N1} & x_{N2} \end{bmatrix}_{N \times 3} = \begin{bmatrix} 1 & x_{11} & 2x_{11} \\ 1 & x_{21} & 2x_{21} \\ \vdots & \vdots & \vdots \\ 1 & x_{N1} & 2x_{N1} \end{bmatrix}_{N \times 3} \quad \left[ \because x_{i2} = 2 \cdot x_{i1} \right]$$

$$X^T = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ x_{11} & x_{21} & x_{31} & \dots & x_{N1} \\ 2x_{11} & 2x_{21} & 2x_{31} & \dots & 2x_{N1} \end{bmatrix}_{3 \times N}$$

$$\therefore X^T X = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ x_{11} & x_{21} & x_{31} & \dots & x_{N1} \\ 2x_{11} & 2x_{21} & 2x_{31} & \dots & 2x_{N1} \end{bmatrix} \begin{bmatrix} 1 & x_{11} & 2x_{11} \\ 1 & x_{21} & 2x_{21} \\ \vdots & \vdots & \vdots \\ 1 & x_{N1} & 2x_{N1} \end{bmatrix}$$

$$= \begin{bmatrix} N & x_{11} + x_{21} + \dots + x_{N1} & 2(x_{11} + x_{21} + \dots + x_{N1}) \\ x_{11} + x_{21} + \dots + x_{N1} & x_{11}^2 + x_{21}^2 + \dots + x_{N1}^2 & 2(x_{11}^2 + x_{21}^2 + \dots + x_{N1}^2) \\ 2(x_{11} + x_{21} + \dots + x_{N1}) & 2(x_{11}^2 + x_{21}^2 + \dots + x_{N1}^2) & 4(x_{11}^2 + x_{21}^2 + \dots + x_{N1}^2) \end{bmatrix}_{3 \times 3}$$


$$= \begin{array}{c} \begin{matrix} C1 & C2 & C3 \end{matrix} \\ \begin{bmatrix} N & x_{11} + x_{21} + \dots + x_{N1} & 2(x_{11} + x_{21} + \dots + x_{N1}) \\ x_{11} + x_{21} + \dots + x_{N1} & x_{11}^2 + x_{21}^2 + \dots + x_{N1}^2 & 2(x_{11}^2 + x_{21}^2 + \dots + x_{N1}^2) \\ 2(x_{11} + x_{21} + \dots + x_{N1}) & 2(x_{11}^2 + x_{21}^2 + \dots + x_{N1}^2) & 4(x_{11}^2 + x_{21}^2 + \dots + x_{N1}^2) \end{bmatrix} \end{array}_{3 \times 3}$$

Since column 2 and column 3 are dependent as  $C3 = 2C2$ , hence the matrix  $X^T X$  is not invertible.  
The rank of the matrix  $(X^T X) = 2$ .

## 5 Problem 6

### 5.1 6a


```
[29]: Laplace = []
Normal = []
x_val = np.arange(-5,5,0.1)
b = 1
for x in x_val:
    Laplace.append((1/2*b) * math.exp(-abs(x)/b))
    Normal.append((1/(b * np.sqrt(2*np.pi)) * np.exp(-(x)**2/(2 * b**2))))
plt.plot(x_val,Laplace, label = "Laplace Distribution")
plt.plot(x_val,Normal, label = "Normal Distribution")
plt.legend()
plt.title("Probability Density Function, for b = 1")
plt.show()
```




```
[30]: Laplace = []
Normal = []
b = 2
for x in x_val:
    Laplace.append((1/2*b) * math.exp(-abs(x)/b))
    Normal.append((1/(b * np.sqrt(2*np.pi)) * np.exp(-(x)**2/(2 * b**2))))
plt.plot(x_val,Laplace, label = "Laplace Distribution")
plt.plot(x_val,Normal, label = "Normal Distribution")
```



```
plt.legend()
plt.title("Probability Density Function, for b = 2")
plt.show()
```



```
[31]: Laplace = []
Normal = []
b = 4
for x in x_val:
    Laplace.append((1/2*b) * math.exp(-abs(x)/b))
    Normal.append((1/(b * np.sqrt(2*np.pi)) * np.exp(-(x)**2/(2 * b**2))))
plt.plot(x_val,Laplace, label = "Laplace Distribution")
plt.plot(x_val,Normal, label = "Normal Distribution")
plt.legend()
plt.title("Probability Density Function, for b = 4")
plt.show()
```



Normal distribution have very short tails whereas the Laplace distribution have longer tails because the Laplace density is expressed in terms of the absolute difference from the mean.

## Problem 6

b)  $y_i = \theta_0 + \theta_1 x_i + \epsilon$

$\epsilon \sim \text{Laplace}(0, b)$  with pdf,  $p(x) = \frac{1}{2b} e^{-|x|/b}$

$$y_i | x_i \sim \text{Laplace}(0, b)$$

$$f(y_i | x_i; \theta, b) = \frac{1}{2b} e^{-|y_i - h_\theta(x_i)|/b} = \frac{1}{2b} e^{-|y_i - h_\theta(x_i)|/b} \quad \rightarrow \text{PDF of the Noise.}$$

Training dataset,

$$f(y_1, \dots, y_N | x_1, \dots, x_N; \theta, b) = \prod_{i=1}^N f(y_i | x_i; \theta, b)$$

The likelihood function,

$$\text{Max } L(\theta) = \prod_{i=1}^N P[y_i | x_i; \theta] = \prod_{i=1}^N f(y_i | x_i; \theta, b)$$

[Assuming, training points are independent]

$$L(\theta) = \prod_{i=1}^N \frac{1}{2b} e^{-|y_i - h_\theta(x_i)|/b}$$

$$= \left(\frac{1}{2b}\right)^N \prod_{i=1}^N e^{-|y_i - h_\theta(x_i)|/b}$$

$$\text{Log } L(\theta) = \underbrace{-N \log(2b)}_{\text{constant}} - \sum_{i=1}^N |y_i - h_\theta(x_i)|/b \rightarrow \text{const.}$$

Maximizing the log likelihood is to minimize the term,

$$\sum_{i=1}^N |y_i - h_\theta(x_i)|$$

$$\therefore J(\theta) = \sum_{i=1}^N |y_i - h_\theta(x_i)|$$

## Problem 6

c)  $J(\theta)$  for Normal Noise.

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N [h_{\theta}(x_i) - y_i]^2 \Rightarrow \text{Error is squared,}$$

$J(\theta)$  for Laplace Noise.

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N |h_{\theta}(x_i) - y_i| \Rightarrow \text{Absolute value of the error}$$

Considering an outlier in the training data, defined as a point of a high residual, the objective with Laplace Noise is more resilient to the effect of outliers. Since the  $J(\theta)$  for Normal Noise is more sensitive to residuals ~~with~~ ~~squar~~ as the residuals are squared; whereas the residuals are not squared in Laplace making it more resilient. ~~The~~